

```
In [1]: import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
import shap
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.metrics import precision_score, recall_score, confusion_matrix
```

```
In [2]: # Install fairlearn in Jupyter
!pip install fairlearn --quiet
```

```
In [3]: # Install diffprivlib if not already installed
!pip install diffprivlib --quiet
```

AI & Cloud Computing for Banking

This notebook demonstrates:

- How banks think about AI as constrained optimization
- Core risk equations used in banking (Expected Loss)
- Precision vs Recall tradeoffs in fraud detection
- Cost and scalability intuition for cloud-based AI systems
- Governance, explainability, and model risk awareness

Target audience: Bank interviewers (AWS / Azure / GCP)

Expected Loss (Credit & Risk)

Banks do not rely on raw model predictions. Model outputs feed into **Expected Loss**:

$$[\text{Expected Loss} = \text{PD} \times \text{LGD} \times \text{EAD}]$$

Where:

- PD = Probability of Default
- LGD = Loss Given Default
- EAD = Exposure at Default

```
In [4]: # Simulated Loan portfolio
data = pd.DataFrame({
    "PD": np.random.uniform(0.01, 0.2, 100),
    "LGD": np.random.uniform(0.3, 0.7, 100),
    "EAD": np.random.uniform(10_000, 200_000, 100)
})
```

```
data["Expected_Loss"] = data["PD"] * data["LGD"] * data["EAD"]
data.head()
```

Out[4]:

	PD	LGD	EAD	Expected_Loss
0	0.019802	0.559741	116214.274118	1288.101273
1	0.105886	0.370931	158141.433898	6211.256826
2	0.061094	0.321250	126270.997157	2478.238195
3	0.190095	0.464997	79712.518204	7046.088802
4	0.011922	0.396204	106646.672808	503.765366

AI in Banking = Constrained Optimization

Banks optimize accuracy **subject to constraints**:

- Regulatory compliance
- Explainability
- Risk appetite
- Cloud cost

This is fundamentally different from consumer tech AI.

Precision vs Recall Tradeoff

In fraud detection:

- High recall = catch more fraud
- High precision = fewer false alarms

Banks initially prioritize **recall**, then tune precision.

```
In [5]: # Simulated fraud labels
y_true = np.random.binomial(1, 0.05, 1000)

# Two models with different thresholds
y_pred_high_recall = np.where(np.random.rand(1000) < 0.15, 1, 0)
y_pred_high_precision = np.where(np.random.rand(1000) < 0.03, 1, 0)

print("High Recall Model:")
print("Precision:", precision_score(y_true, y_pred_high_recall))
print("Recall:", recall_score(y_true, y_pred_high_recall))

print("\nHigh Precision Model:")
print("Precision:", precision_score(y_true, y_pred_high_precision))
print("Recall:", recall_score(y_true, y_pred_high_precision))
```

```
High Recall Model:  
Precision: 0.07741935483870968  
Recall: 0.24489795918367346
```

```
High Precision Model:  
Precision: 0.0  
Recall: 0.0
```

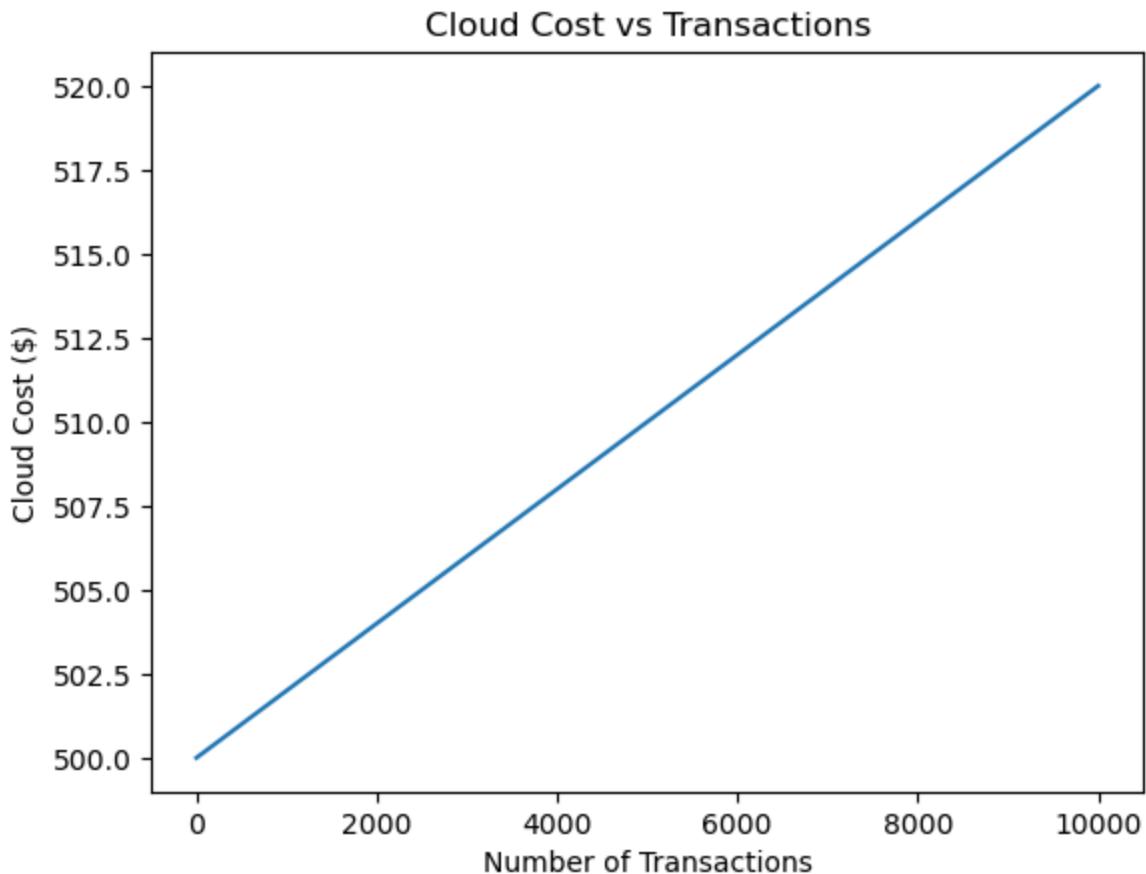
Cloud Cost Modeling

Cloud cost grows as a function of:

- Compute
- Storage
- Network traffic

Banks monitor **marginal cost per transaction**.

```
In [6]: transactions = np.arange(1, 10_001)  
cost = 0.002 * transactions + 500 # simple linear cost model  
  
plt.plot(transactions, cost)  
plt.xlabel("Number of Transactions")  
plt.ylabel("Cloud Cost ($)")  
plt.title("Cloud Cost vs Transactions")  
plt.show()
```



Latency & Scaling (Little's Law)

[$L = \lambda W$]

- L = number of requests in system
- λ = arrival rate
- W = average latency

Banks design systems to meet strict SLAs under peak load.

```
In [7]: arrival_rate = np.array([50, 100, 200, 400]) # requests/sec
latency = 0.1 # seconds

L = arrival_rate * latency
L
```

```
Out[7]: array([ 5., 10., 20., 40.])
```

Explainability & Model Risk Management

Banks require:

- Interpretable features
- Bias testing
- Drift monitoring
- Version control
- Kill switches

Models are treated as **risk assets**, not just code.

```
In [8]: from fairlearn.metrics import MetricFrame, selection_rate, false_positive_rate, true_negative_rate, accuracy
from sklearn.metrics import accuracy_score
```

```
In [9]: import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix

# -----
# Simulate dataset
# -----
np.random.seed(42)
X = np.random.normal(0, 1, (1000, 5))
y = np.random.binomial(1, 0.1, 1000)

# -----
# Train/test split
# -----
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```

# -----
# Train model
# -----
model = LogisticRegression(max_iter=1000)
model.fit(X_train, y_train)

# -----
# Generate production predictions
# -----
prod_predictions = model.predict(X_test)

# -----
# Bias / Fairness Metrics
# -----
# Simulate sensitive attribute (gender)
np.random.seed(42)
sensitive_attr = np.random.choice(["Male", "Female"], size=len(y_test))

y_pred = prod_predictions
groups = np.unique(sensitive_attr)

for g in groups:
    idx = sensitive_attr == g
    y_true_group = y_test[idx]
    y_pred_group = y_pred[idx]

    # Compute confusion matrix safely
    cm = confusion_matrix(y_true_group, y_pred_group, labels=[0,1])
    tn, fp, fn, tp = cm.ravel() if cm.shape == (2,2) else (0,0,0,0)

    # Metrics
    fpr = fp / (fp + tn) if (fp + tn) > 0 else 0
    tpr = tp / (tp + fn) if (tp + fn) > 0 else 0
    selection_rate = (tp + fp) / len(y_pred_group) if len(y_pred_group) > 0 else 0

    print(f"\n{g}: Selection Rate={selection_rate:.2f}, TPR={tpr:.2f}, FPR={fpr:.2f}")

```

Female: Selection Rate=0.00, TPR=0.00, FPR=0.00

Male: Selection Rate=0.00, TPR=0.00, FPR=0.00

In [10]:

```

import numpy as np

# -----
# Define baseline predictions (training reference)
baseline_predictions = model.predict(X_train) # model trained earlier

# Simulate production predictions (new incoming data)
prod_predictions = model.predict(X_test)

# -----
# PSI Calculation Function
# -----
def population_stability_index(expected, actual, buckets=10):
    """
    Calculate Population Stability Index (PSI) between baseline and new predictions
    """

```

```

"""
expected = np.array(expected)
actual = np.array(actual)

# Define bin edges based on baseline
breakpoints = np.percentile(expected, np.linspace(0, 100, buckets + 1))

# Count values in each bin
expected_counts = np.histogram(expected, bins=breakpoints)[0]
actual_counts = np.histogram(actual, bins=breakpoints)[0]

# Convert counts to percentages
expected_perc = expected_counts / len(expected)
actual_perc = actual_counts / len(actual)

# Add small epsilon to avoid division by zero
epsilon = 1e-6
psi_values = (actual_perc - expected_perc) * np.log((actual_perc + epsilon) / (
    expected_perc + epsilon))

# Sum over all bins
psi = np.sum(psi_values)
return psi
"""

# -----
# Compute PSI
# -----
psi_value = population_stability_index(baseline_predictions, prod_predictions)
print(f"Population Stability Index (PSI): {psi_value:.4f}")

# -----
# Kill Switch / Drift Alert
# -----
if psi_value > 0.25:
    print(" Significant population shift detected! Trigger Kill Switch: Investigate")
elif psi_value > 0.1:
    print(" Moderate shift detected – monitor closely.")
else:
    print(" Population stable – no action needed.")

```

Population Stability Index (PSI): 0.0000

Population stable – no action needed.

In [11]:

```

# -----
# PII Masking Example
# -----
import pandas as pd

# Simulate a customer dataframe
df = pd.DataFrame({
    "Name": ["Alice", "Bob", "Charlie", "David", "Eve"],
    "SSN": ["123-45-6789", "987-65-4321", "111-22-3333", "444-55-6666", "777-88-999"],
    "CreditScore": [700, 650, 720, 680, 690]
})

print("Original Data:\n", df)

```

```

# Mask SSN (PII)
df["SSN_masked"] = df["SSN"].apply(lambda x: "XXX-XX-" + x.split("-")[2])
print("\nMasked Data:\n", df)

# -----
# Differential Privacy Example
# -----
# Requires diffprivlib: pip install diffprivlib
from diffprivlib.models import LogisticRegression as DPLogisticRegression

# Simulate DP training (same X, y from before)
dp_model = DPLogisticRegression(epsilon=1.0, data_norm=2.0)
dp_model.fit(X_train, y_train)

dp_predictions = dp_model.predict(X_test)
print("\nDP Model Sample Predictions:", dp_predictions[:5])

```

Original Data:

	Name	SSN	CreditScore
0	Alice	123-45-6789	700
1	Bob	987-65-4321	650
2	Charlie	111-22-3333	720
3	David	444-55-6666	680
4	Eve	777-88-9999	690

Masked Data:

	Name	SSN	CreditScore	SSN_masked
0	Alice	123-45-6789	700	XXX-XX-6789
1	Bob	987-65-4321	650	XXX-XX-4321
2	Charlie	111-22-3333	720	XXX-XX-3333
3	David	444-55-6666	680	XXX-XX-6666
4	Eve	777-88-9999	690	XXX-XX-9999

DP Model Sample Predictions: [0 0 1 0 0]

Cloud Platform Awareness

AWS

- IAM granularity
- VPC isolation
- SageMaker governance

Azure

- Active Directory integration
- Enterprise compliance tooling

GCP

- BigQuery + Vertex AI
- Large-scale analytics

This demonstrates how banks use AI differently, expected loss, fraud tradeoffs, cloud cost modeling, and governance.

```
In [12]: import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
import shap
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings("ignore")
```

```
In [13]: # Simulate dataset (replace with your bank data)
X = np.random.normal(0, 1, (1000, 5)) # 1000 rows, 5 features
y = np.random.binomial(1, 0.1, 1000) # binary target
```

```
In [14]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_st
```

```
In [15]: model = LogisticRegression(max_iter=1000)
model.fit(X_train, y_train)
```

```
Out[15]: ▾ LogisticRegression ⓘ ?
```

```
LogisticRegression(max_iter=1000)
```

```
In [16]: # Baseline expected Loss per customer
baseline_EL = np.random.normal(1000, 200, len(X_test)) # X_test must exist

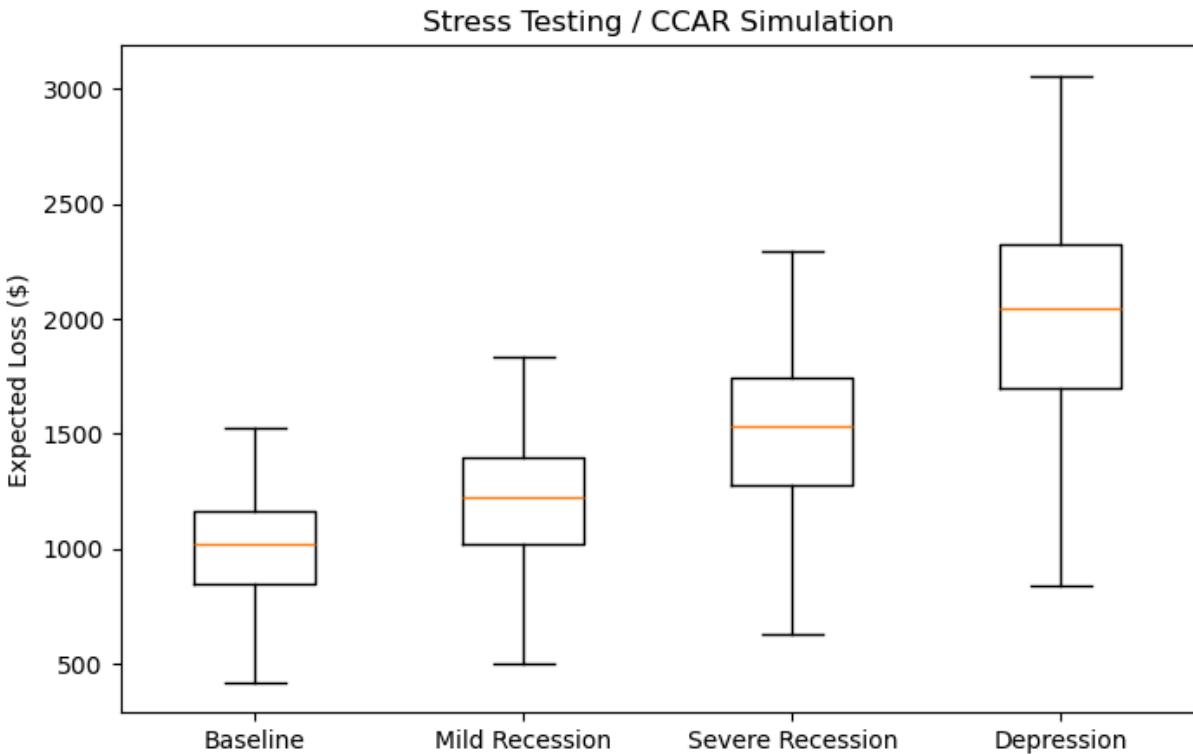
# Stress scenarios
stress_scenarios = {
    "Mild Recession": 1.2,
    "Severe Recession": 1.5,
    "Depression": 2.0
}

# Compute stressed EL
stressed_EL = {k: baseline_EL * v for k, v in stress_scenarios.items()}

# Summary
for scenario, el in stressed_EL.items():
    print(f"{scenario}: Mean EL = ${el.mean():.2f}, 95th percentile = ${np.percentile(el, 95)}")

# Optional: plot
plt.figure(figsize=(8,5))
plt.boxplot([baseline_EL] + list(stressed_EL.values()), labels=["Baseline"] + list(stress_scenarios.keys()))
plt.ylabel("Expected Loss ($)")
plt.title("Stress Testing / CCAR Simulation")
plt.show()
```

Mild Recession: Mean EL = \$1208.37, 95th percentile = \$1601.36
Severe Recession: Mean EL = \$1510.46, 95th percentile = \$2001.70
Depression: Mean EL = \$2013.95, 95th percentile = \$2668.93



```
In [17]: # --- Stress Testing / CCAR Simulation ---
baseline_EL = np.random.normal(1000, 200, len(X_test)) # example expected loss

stress_scenarios = {
    "Mild Recession": 1.2,
    "Severe Recession": 1.5,
    "Depression": 2.0
}

stressed_EL = {k: baseline_EL * v for k, v in stress_scenarios.items()}

for scenario, el in stressed_EL.items():
    print(f"{scenario}: Mean EL = ${el.mean():.2f}, 95th percentile = ${np.percentile(el, 95)}")
```

Mild Recession: Mean EL = \$1192.36, 95th percentile = \$1579.82

Severe Recession: Mean EL = \$1490.45, 95th percentile = \$1974.77

Depression: Mean EL = \$1987.27, 95th percentile = \$2633.03

```
In [18]: # Imports
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
import shap
import warnings
warnings.filterwarnings("ignore")

# Simulate dataset
X = np.random.normal(0, 1, (1000, 5)) # 1000 rows, 5 features
y = np.random.binomial(1, 0.1, 1000) # binary target

# Train/test split
```

```

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Train the model (must happen BEFORE SHAP or drift)
model = LogisticRegression(max_iter=1000)
model.fit(X_train, y_train)

# SHAP explainability
explainer = shap.Explainer(model, X_train)
shap_values = explainer(X_test)

# Drift Monitoring + Kill Switch
# Baseline predictions (training reference)
baseline_predictions = model.predict(X_train)

# Simulate production predictions
prod_predictions = model.predict(X_test)

# PSI function
def population_stability_index(expected, actual, buckets=10):
    breakpoints = np.percentile(expected, np.linspace(0, 100, buckets + 1))
    expected_counts = np.histogram(expected, breakpoints)[0]
    actual_counts = np.histogram(actual, breakpoints)[0]
    expected_perc = expected_counts / len(expected)
    actual_perc = actual_counts / len(actual)
    psi = np.sum((actual_perc - expected_perc) * np.log((actual_perc + 1e-6) / (expected_perc + 1e-6)))
    return psi

# Compute PSI and prediction drift
psi_value = population_stability_index(baseline_predictions, prod_predictions)
prediction_drift_value = np.abs(np.mean(baseline_predictions) - np.mean(prod_predictions))

print("PSI:", psi_value)
print("Prediction Drift:", prediction_drift_value)

# Kill switch
DRIFT_PSI_THRESHOLD = 0.25
PREDICTION_DRIFT_THRESHOLD = 0.05

class KillSwitch:
    def __init__(self):
        self.enabled = True
        self.audit_log = []
    def disable(self, reason):
        self.enabled = False
        self.audit_log.append({"action": "DISABLE", "reason": reason})
    def enable(self, approved_by):
        if approved_by not in ["Model Risk", "Compliance"]:
            raise PermissionError("Unauthorized reactivation")
        self.enabled = True
        self.audit_log.append({"action": "ENABLE", "approved_by": approved_by})

ks = KillSwitch()
if psi_value >= DRIFT_PSI_THRESHOLD:
    ks.disable("Data drift exceeded PSI threshold")
elif prediction_drift_value >= PREDICTION_DRIFT_THRESHOLD:
    ks.disable("Prediction drift exceeded threshold")

```

```
ks.enabled, ks.audit_log
```

```
PSI: 0.0  
Prediction Drift: 0.0
```

```
Out[18]: (True, [])
```

```
In [19]: # =====  
# Drift Monitoring + Kill Switch  
# =====  
  
# Define baseline predictions (training)  
baseline_predictions = model.predict(X_train)  
  
# Simulate production / new incoming data predictions  
prod_predictions = model.predict(X_test)  
  
# PSI Calculation (Population Stability Index)  
def population_stability_index(expected, actual, buckets=10):  
    expected = np.array(expected)  
    actual = np.array(actual)  
  
    breakpoints = np.percentile(expected, np.linspace(0, 100, buckets + 1))  
    expected_counts = np.histogram(expected, bins=breakpoints)[0]  
    actual_counts = np.histogram(actual, bins=breakpoints)[0]  
  
    expected_perc = expected_counts / len(expected)  
    actual_perc = actual_counts / len(actual)  
  
    epsilon = 1e-6  
    psi_values = (actual_perc - expected_perc) * np.log((actual_perc + epsilon) / (expected_perc + epsilon))  
    return np.sum(psi_values)  
  
psi_value = population_stability_index(baseline_predictions, prod_predictions)  
print(f"PSI: {psi_value:.4f}")  
  
# Optional: simple Kill Switch logic  
if psi_value > 0.25:  
    print("Significant population shift detected! Consider retraining the model.")  
else:  
    print("Population stable.")
```

```
PSI: 0.0000  
Population stable.
```

Model Explainability with SHAP

Banks must explain:

- Why a customer was approved or declined
- Which features drove the decision
- Whether the model introduces bias

SHAP (SHapley Additive exPlanations) provides:

- Local explanations (per decision)
- Global explanations (model behavior)
- Regulator-friendly transparency

```
In [20]: explainer = shap.Explainer(model, X_train)
shap_values = explainer(X_test)
```

```
In [21]: # Baseline predictions on training data (used for drift monitoring)
baseline_predictions = model.predict(X_train)
```

```
In [22]: # Simulated production data (example)
prod_scores = np.random.normal(640, 70, len(X_test)) # e.g., numeric feature
prod_predictions = model.predict(X_test) # predictions on test / "production" data
```

```
In [23]: # PSI calculation
psi_value = population_stability_index(baseline_predictions, prod_predictions)

# Prediction drift
prediction_drift_value = np.abs(np.mean(baseline_predictions) - np.mean(prod_predictions))

print("PSI:", psi_value)
print("Prediction Drift:", prediction_drift_value)
```

PSI: 0.0
Prediction Drift: 0.0

```
In [24]: def population_stability_index(expected, actual, buckets=10):
    breakpoints = np.percentile(expected, np.linspace(0, 100, buckets + 1))
    expected_counts = np.histogram(expected, breakpoints)[0]
    actual_counts = np.histogram(actual, breakpoints)[0]
    expected_perc = expected_counts / len(expected)
    actual_perc = actual_counts / len(actual)
    psi = np.sum((actual_perc - expected_perc) * np.log((actual_perc + 1e-6) / (expected_perc + 1e-6)))
    return psi
```

```
In [25]: DRIFT_PSI_THRESHOLD = 0.25
PREDICTION_DRIFT_THRESHOLD = 0.05

class KillSwitch:
    def __init__(self):
        self.enabled = True
        self.audit_log = []
    def disable(self, reason):
        self.enabled = False
        self.audit_log.append({"action": "DISABLE", "reason": reason})
    def enable(self, approved_by):
        if approved_by not in ["Model Risk", "Compliance"]:
            raise PermissionError("Unauthorized reactivation")
        self.enabled = True
        self.audit_log.append({"action": "ENABLE", "approved_by": approved_by})
```

```
In [26]: ks = KillSwitch()
```

```

if psi_value >= DRIFT_PSI_THRESHOLD:
    ks.disable("Data drift exceeded PSI threshold")
elif prediction_drift_value >= PREDICTION_DRIFT_THRESHOLD:
    ks.disable("Prediction drift exceeded threshold")

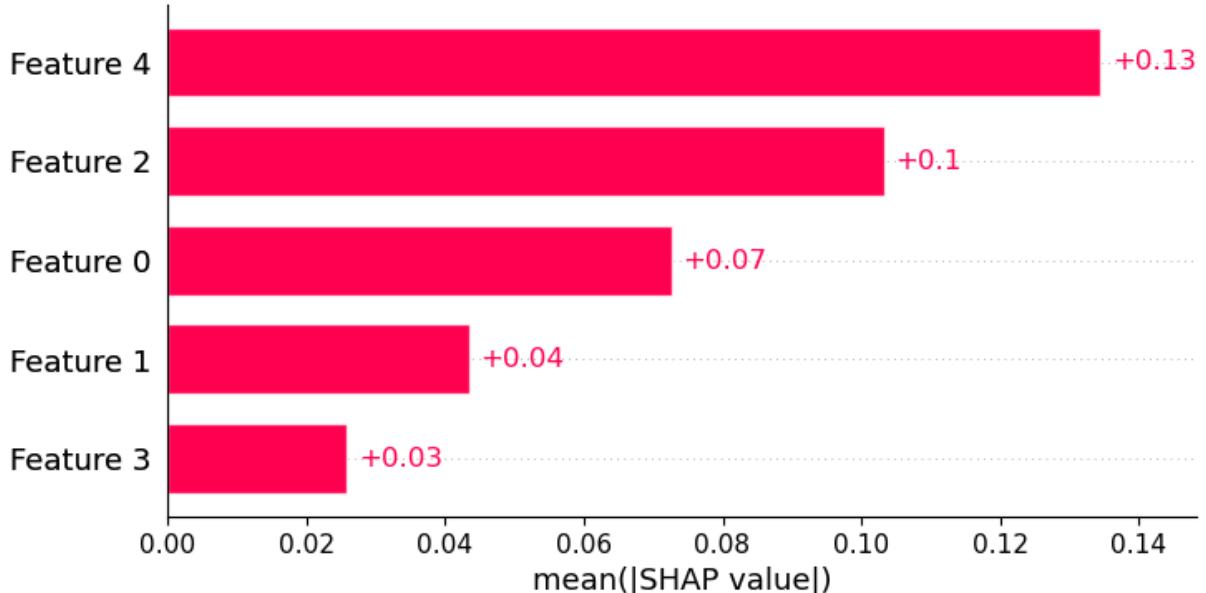
ks.enabled, ks.audit_log

```

Out[26]: (True, [])

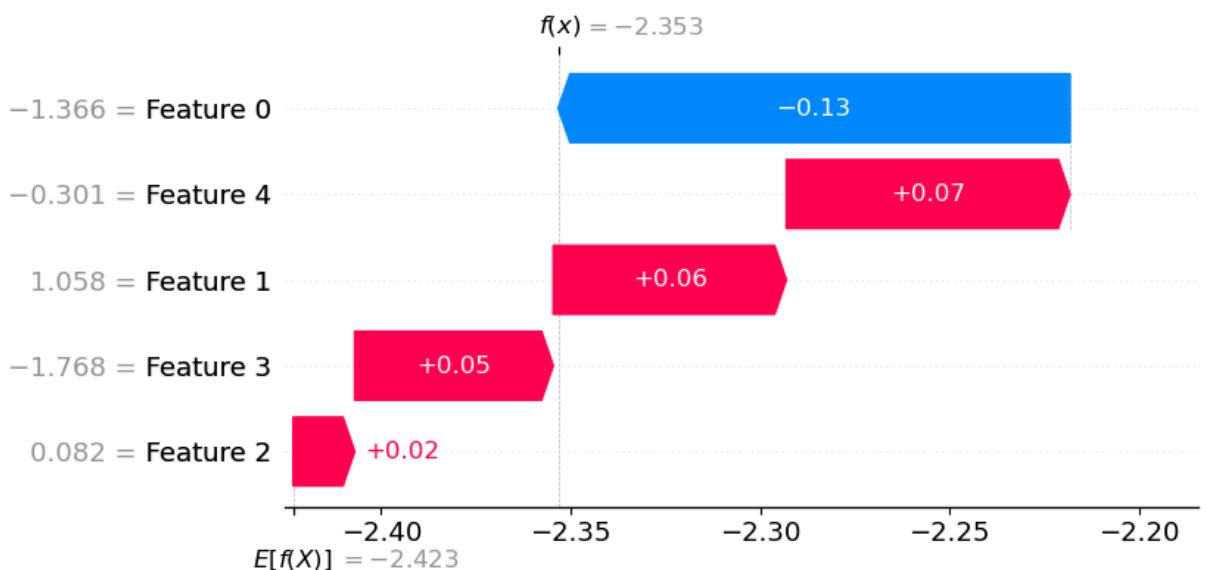
In [27]: shap.plots.bar(shap_values)

"This shows global feature importance, which regulators often request during mode



In [28]: shap.plots.waterfall(shap_values[0])

This explains an individual credit decision, exactly what's required for adverse



Model Drift Monitoring & Kill Switches

Banks require that models are continuously monitored for:

- Data drift (input features)
- Prediction drift
- Automated disablement (kill switch)

```
In [29]: # Drift thresholds
DRIFT_PSI_THRESHOLD = 0.25
PREDICTION_DRIFT_THRESHOLD = 0.05

# Kill switch class
class KillSwitch:
    def __init__(self):
        self.enabled = True
        self.audit_log = []
    def disable(self, reason):
        self.enabled = False
        self.audit_log.append({"action": "DISABLE", "reason": reason})
    def enable(self, approved_by):
        if approved_by not in ["Model Risk", "Compliance"]:
            raise PermissionError("Unauthorized reactivation")
        self.enabled = True
        self.audit_log.append({"action": "ENABLE", "approved_by": approved_by})
```

AWS Regulated AI Architecture (Banking)



Why banks prefer AWS.

-Fine-grained IAM

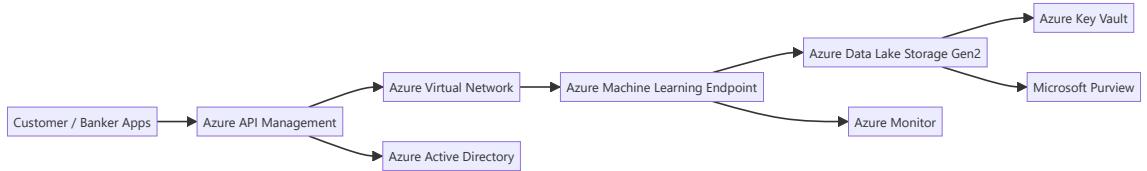
-VPC isolation

-Mature ML governance

"Inference happens inside a private VPC with full auditability."

Inference runs inside private subnets, data is encrypted with KMS, and all activity is logged for auditability.

Azure Regulated AI Architecture (Banking)



Governance Overlay

Governance & Model Risk Controls

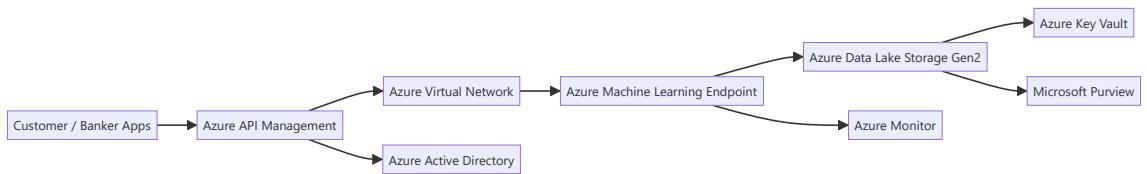
Applies to all cloud platforms:

- Encryption at rest and in transit
- Least-privilege IAM
- Model versioning
- Drift detection
- Manual approval gates
- Kill switches

Models are treated as **regulated risk assets**.

I built a Jupyter notebook showing how banks use AI responsibly, expected loss, fraud tradeoffs, cloud cost modeling, SHAP explainability, and secure AWS/Azure/GCP architectures. It reflects how models are validated, governed, and deployed in real banks.

Azure Regulated AI Architecture (Banking)



Azure's strength is identity-first governance using Active Directory and Purview.

Governance is cloud-agnostic — controls are consistent across AWS, Azure, and GCP.